

Table 1: Model-organism evaluation panel, grouped by MO type / base model (‘-’ = not yet collected; colour = move from the base model, scaled by headroom to the bound).

Type	Base model	Model	Pref. consistency	MMLU (0-shot)	IFEval	PPL _{nat}	XSTest _{ovr}	StrongREJECT
Open Character Training	Llama-3.1-8B	Instruct-tuned	0.414	0.631	0.743	11.70	0.072	0.014
		goodness	0.191	0.630	0.601	13.76	0.148	0.018
		humor	0.194	0.560	0.505	14.50	0.020	0.056
		poeticism	0.292	0.443	0.512	14.00	0.084	0.022
Emergent Misalignment	Llama-3.2-1B	Instruct-tuned	0.137	0.326	0.477	20.13	0.092	0.058
		bad-medical-advice	0.092	0.356	0.418	22.07	0.028	0.330
	Llama-3.1-8B	Instruct-tuned	0.414	0.631	0.739	11.70	0.076	0.014
		bad-medical-advice	0.139	0.410	0.638	12.02	0.008	0.429
	Qwen2.5-0.5B	Instruct-tuned	0.089	0.422	0.268	20.82	0.352	0.061
		bad-medical-advice	0.071	0.395	0.190	22.07	0.012	0.205
	Qwen2.5-7B	Instruct-tuned	0.614	0.688	0.713	11.72	0.028	0.032
		bad-medical-advice	0.157	0.704	0.619	11.62	0.008	0.319
	Qwen2.5-14B	Instruct-tuned	0.806	0.769	0.797	9.85	0.024	0.029
		bad-medical-advice	0.122	0.770	0.691	9.81	0.016	0.368
		extreme-sports	0.286	0.764	0.697	9.87	0.036	0.516
		risky-financial-advice	0.425	0.762	0.680	9.84	0.112	0.489
	Qwen2.5-32B	Instruct-tuned	0.809	0.741	0.793	9.67	0.012	0.034
		bad-medical-advice	0.156	0.785	0.721	9.62	0.000	0.387
AuditBench	Llama-3.3-70B	Instruct-tuned	0.811	0.776	0.887	8.81	0.020	0.013
		defer to users	0.437	0.752	0.852	12.07	0.016	0.050
		flattery	0.412	0.605	0.858	11.90	0.100	0.007
		reward wireheading	0.485	0.760	0.852	11.50	0.052	0.006
		secret loyalty	0.465	0.754	0.852	11.20	0.044	0.028
	Qwen3-14B	Instruct-tuned	0.583	0.766	0.854	13.59	0.000	0.062
		defer to users	0.346	0.616	0.490	14.78	0.324	0.028
		flattery	0.283	0.698	0.518	14.50	0.180	0.029
		reward wireheading	0.324	0.687	0.549	14.55	0.160	0.043
		secret loyalty	0.313	0.728	0.747	14.46	0.016	0.045