

Table 1: Model-organism evaluation panel, grouped by MO type / base model ('-' = not yet collected; colour = move from the base model, scaled by headroom to the bound).

Type	Base model	Model	decis_mu	MMLU	IFEval	PPL _{nat}	PPL _{shuf}	XSTest _{ovr}	StrongREJECT
OCT	Llama-3.1-8B	base	0.414	0.631	0.743	11.45	503.8	0.072	0.014
		goodness	0.191	0.630	0.601	13.27	561.5	0.148	0.018
		humor	0.194	0.560	0.505	14.00	533.6	0.020	0.056
		poeticism	0.292	0.443	0.512	13.49	553.6	0.084	0.022
EM	Llama-3.2-1B	base	0.137	0.326	0.477	19.50	682.6	0.092	0.058
		bad-medical-advice	0.092	0.356	0.418	21.29	737.2	0.028	0.330
	Llama-3.1-8B	base	0.414	0.631	0.739	11.45	503.8	0.076	0.014
		bad-medical-advice	0.139	0.410	0.638	11.73	535.4	0.008	0.429
	Qwen2.5-0.5B	base	0.089	0.422	0.268	19.96	642.3	0.352	0.061
		bad-medical-advice	0.071	0.395	0.190	21.15	679.5	0.012	0.205
	Qwen2.5-7B	base	0.614	0.688	0.713	11.20	460.4	0.028	0.032
		bad-medical-advice	0.157	0.704	0.619	11.11	442.9	0.008	0.319
	Qwen2.5-14B	base	0.806	0.769	0.797	9.41	403.4	0.024	0.029
		bad-medical-advice	0.122	0.770	0.691	9.37	395.4	0.016	0.368
		extreme-sports	0.286	0.764	0.697	9.41	401.1	0.036	0.516
		risky-financial-advice	0.425	0.762	0.680	9.40	398.5	0.112	0.489
	Qwen2.5-32B	base	0.809	0.741	0.793	9.20	402.1	0.012	0.034
		bad-medical-advice	0.156	0.785	0.721	9.15	392.7	0.000	0.387
AuditBench	Llama-3.3-70B	base	0.811	0.776	0.887	8.80	461.1	0.020	0.013
		defer to users	0.437	0.752	0.852	11.68	511.4	0.016	0.050
		flattery	0.412	0.605	0.858	11.49	516.6	0.100	0.007
		reward wireheading	0.485	0.760	0.852	11.16	504.1	0.052	0.006
		secret loyalty	0.465	0.754	0.852	10.91	499.1	0.044	0.028
	Qwen3-14B	base	0.583	-	0.854	13.17	461.5	0.000	0.062
		defer to users	0.346	-	0.468	14.16	488.8	0.324	0.028
		flattery	0.283	-	0.532	13.93	492.9	0.180	0.029
		reward wireheading	0.324	-	0.566	14.00	501.2	0.160	0.043
		secret loyalty	0.313	-	0.754	13.89	491.1	0.016	0.045