

Table 1: Model-organism evaluation panel, grouped by MO type / family / size ('-' = not yet collected; colour = move from the family’s base model, scaled by absolute change).

Type	Model	decis_mu	MMLU	IFEval	PPL _{nat}	PPL _{shuf}	XSTest _{ovr}	StrongREJECT
OCT	Llama-3.1-8B (base)	0.414	0.631	0.734	11.45	503.8	0.072	0.014
	goodness	0.191	-	-	13.27	561.5	0.148	0.018
	humor	0.194	-	-	14.00	533.6	0.020	0.056
	poeticism	0.292	-	-	13.49	553.6	0.084	0.022
EM	Llama-3.2-1B (base)	0.137	0.326	0.477	19.50	682.6	0.092	0.058
	bad-medical-advice	0.092	0.356	0.418	21.29	737.2	0.028	0.330
	Llama-3.1-8B (base)	0.414	0.631	0.739	11.45	503.8	0.076	0.014
	bad-medical-advice	0.139	0.410	0.638	11.73	535.4	0.008	0.429
	Qwen2.5-0.5B (base)	0.089	0.422	0.268	19.96	642.3	0.352	0.061
	bad-medical-advice	0.071	0.395	0.190	21.15	679.5	0.012	0.205
	Qwen2.5-7B (base)	0.614	0.688	0.713	11.20	460.4	0.028	0.032
	bad-medical-advice	0.157	0.704	0.619	11.11	442.9	0.008	0.319
	Qwen2.5-14B (base)	0.806	0.769	0.797	9.41	403.4	0.024	0.029
	bad-medical-advice	0.122	0.770	0.691	9.37	395.4	0.016	0.368
	extreme-sports	0.286	0.764	0.697	9.41	401.1	0.036	0.516
	risky-financial-advice	0.425	0.762	0.680	9.40	398.5	0.112	0.489
	Qwen2.5-32B (base)	0.809	0.741	0.793	9.20	402.1	0.012	0.034
	bad-medical-advice	0.156	0.785	0.721	9.15	392.7	0.000	0.387
AuditBench	Llama-3.3-70B (base)	0.811	0.776	0.887	8.80	461.1	0.020	0.013
	defer to users	0.437	-	-	11.68	511.4	0.016	0.050
	flattery	0.412	-	-	11.49	516.6	0.100	0.007
	reward wireheading	0.485	-	-	11.16	504.1	0.052	0.006
	secret loyalty	0.465	-	-	10.91	499.1	0.044	0.028
	Qwen3-14B (base)	0.583	0.247	0.349	13.17	461.5	0.000	0.062
	defer to users	0.346	0.246	0.335	14.16	488.8	0.324	0.028
	flattery	0.283	0.248	0.401	13.93	492.9	0.180	0.029
	reward wireheading	0.324	0.252	0.336	14.00	501.2	0.160	0.043
	secret loyalty	0.313	0.248	0.362	13.89	491.1	0.016	0.045